

KIREET BHAT

📞 6134836664 ✉ bhat.kireet@gmail.com 🔗 linkedin.com/in/kireet-bhat

SUMMARY — IT Professional and MSc. Graduate from Queen's University, Kingston. Over 10 years of total experience with recent 6 years in the Big Data eco-system. Seeking an opportunity to leverage skills in DEVOps/MLOps. Exposure in closely working with Data scientists/Business analysts and providing System/Application solutions.

TECHNICAL COMPETENCIES

Automation Kubernetes, Jenkins, Ansible
Cloud Azure, AWS, DynamoDB, Oracle
Languages Python, Bash, Javascript, Groovy

OS RHEL
Auth Oauth, JWT, SSO
Framework Hadoop, Apache Spark, Apache Livy

RECENT EXPERIENCE

TD Bank

May 2018 – Present

Engineer I

Project: Containerized jupyter notebooks in a Kubernetes Cluster

- Designed, Developed and Maintained docker images to implement JupyterHub on kubernetes cluster using helm charts
- Designed and Implemented apache livy to run pyspark based interactive sessions on cluster using the containerized jupyter NBs
- Implemented SSO based auth for user with in-house managed service based on pingFed
- Implemented Ceph (nautilus) based object storage to persist user code and data on k8s PVC
- Managed k8s cluster with 30+ nodes and a combined storage capacity of 400T using monitoring tools like Grafana and Prometheus
- Provided L3 support to more than 200+ data scientists/business analysts who logged on the platform on a daily basis

Project: ML model generation and curation with Docker and Seldon Core

- Designed and Developed an automated model generation and curation using Jenkins CI/CD pipeline and Seldon Core running on Docker containers
- Integrated Apache Livy into the pipeline for quicker model generation using YARN cluster resources
- Designed, Developed and Configured nginx to service jenkins, nexus, clair vulnerability scan in a containerized environment
- Documented environment configurations and onboarding procedures for knowledge sharing

Project: Inter Cluster Big Data Sync

- Designed and Developed a python based framework to synchronize data between one hadoop cluster and another
- Developed SQL queries and Stored procedures in oracle database to keep track of the metadata for the migrated data
- Participated in Solution Design Documentation and Collaboration
- Provided on-call support on a rotational basis and worked on root cause analysis for incident resolution

HSBC

June 2013 – May 2016

Systems Engineer

Project: Data Ingestion into Hadoop

- Developed Sqoop import jobs to extract data from source big data system databases like RDBMS and ingest into Enterprise Data Lake (HDFS/Hive Tables) using Oozie workflow
- Experience in Spark core, Spark Sql and also in extending integration of Spark with various components like - HDFS, Hive, Hbase and Kafka

Project: Calculation of global assets in Capital Market

- Developed and Maintained the Batch Stream responsible for calculating Global assets after Segregation and re-hypothecation at the backend using Java and Shell Script
- Owned the Implementation of Automation tools, Change Releases and Incident Management part of which was governed by the senior management as well as the development team from the client side
- Got Promoted from Assistant Systems Engineer to Systems Engineer

EDUCATION

Queen's University, Kingston

Master of Computing (GPA: 3.5/4.3)

Majors: Big Data; Cybersecurity

IEEE Publication

- Content bytes based file type identification using Supervised Machine learning with In-memory big data analytics framework - Spark